

Computational Statistics & Probability

Problem Set 3 - Information Criteria and Interactions

Author: Neelesh Bhalla

Collaborators: Nils Marthiensen, Chia-Jung Chang

2022-11-30

1. Collider Bias and Information Criteria

Return to the textbook example in §6.3.1, which explores the relationship between age, marriage and happiness.


```
library(rethinking)

## Loading required package: rstan
## Loading required package: StanHeaders
##
## rstan version 2.26.13 (Stan version 2.26.1)
## For execution on a local, multicore CPU with excess RAM we recommend calling
## options(mc.cores = parallel::detectCores()).
## To avoid recompilation of unchanged Stan programs, we recommend calling
## rstan_options(auto_write = TRUE)
## For within-chain threading using `reduce_sum()` or `map_rect()` Stan functions,
## change `threads_per_chain` option:
## rstan_options(threads_per_chain = 1)
## Loading required package: cmdstanr
## This is cmdstanr version 0.5.3
## - CmdStanR documentation and vignettes: mc-stan.org/cmdstanr
## - CmdStan path: /Users/neelesh/.cmdstan/cmdstan-2.30.1
## - CmdStan version: 2.30.1
##
## A newer version of CmdStan is available. See ?install_cmdstan() to install it.
## To disable this check set option or environment variable CMDSTANR_NO_VER_CHECK=TRUE.
## Loading required package: parallel
## rethinking (Version 2.21)
##
## Attaching package: 'rethinking'
## The following object is masked from 'package:rstan':
##
##      stan
```

```
## The following object is masked from 'package:stats':
##
##      rstudent
d <- sim_happiness( seed=1515 , N_years=1000)
d2 <- d[ d$age>17 , ] # only adults
d2$A <- (d2$age - 18) / (65 - 18)
d2$mid <- d2$married + 1
precis(d)[,1:4]

##              mean      sd      5.5%      94.5%
## age      3.300000e+01 18.7688832  4.000000 62.000000
## married   2.930769e-01  0.4553486  0.000000  1.000000
## happiness -1.000070e-16  1.2144211 -1.789474  1.789474

# drawing the D.A.G.
library(dagitty)
dag_q1 <- dagitty('dag{ H -> M <- A }')
drawdag( dag_q1 )
```



a) Which model is expected to make better predictions according to these information criteria?

```
# recalling model m6.9 that considers the effect of age and marriage status on happiness
m6.9 <- quap(
  alist(
    happiness ~ dnorm( mu , sigma ),
    mu <- a[mid] + bA*A,
    a[mid] ~ dnorm( 0 , 1 ),
    bA ~ dnorm( 0 , 2 ),
    sigma ~ dexp(1)
  ) , data=d2 )
```

```
precis(m6.9,depth=2)
```

```
##              mean          sd        5.5%        94.5%
## a[1]   -0.1947491  0.06521259 -0.2989714 -0.0905268
## a[2]    1.2161249  0.08876809  1.0742563  1.3579934
## bA     -0.7325968  0.11708179 -0.9197161 -0.5454775
## sigma   1.0199308  0.02325838  0.9827594  1.0571022
```

model m6.9 is quite sure that age is negatively associated with happiness.

recalling model m6.10 that omits marriage status

```
m6.10 <- quap(
  alist(
    happiness ~ dnorm( mu , sigma ),
    mu <- a + bA*A,
    a ~ dnorm( 0 , 1 ),
    bA ~ dnorm( 0 , 2 ),
    sigma ~ dexp(1)
  ) , data=d2 )
precis(m6.10)
```

```
##              mean          sd        5.5%        94.5%
## a      -9.798009e-06  0.07674935 -0.1226701  0.1226505
## bA      3.917361e-05  0.13225839 -0.2113353  0.2114136
## sigma   1.213175e+00  0.02766008  1.1689688  1.2573811
```

model m6.10 in contrast, finds no association between age and happiness.

Notes from the book (pg 184):

*# The pattern above is exactly what we should expect when we condition on a collider.
 # The collider is marriage status. It a common consequence of age and happiness.
 # As a result, when we condition on it, we induce a spurious association between the two causes.
 # So it looks like, to model m6.9, that age is negatively associated with happiness.
 # But this is just a statistical association, not a causal association.
 # Once we know whether someone is married or not, then their age does provide information
 # about how happy they are.*

Compare the two models, m6.9 and m6.10, using both PSIS and WAIC.

comparing models using PSIS

```
compare( m6.9, m6.10, func=PSIS)
```

```
##              PSIS          SE    dPSIS      dSE    pPSIS      weight
## m6.9   2771.068  37.13856    0.0000      NA  3.752925  1.000000e+00
## m6.10  3102.028  27.71311  330.9597  34.1829  2.415411  1.358367e-72
```

comparing models using WAIC

```
compare( m6.9, m6.10, func=WAIC)
```

```
##              WAIC          SE    dWAIC      dSE    pWAIC      weight
## m6.9   2771.393  36.93196    0.0000      NA  3.960565  1.000000e+00
## m6.10  3101.906  27.68192  330.5132  34.02321  2.330043  1.698146e-72
```

a) Which model is expected to make better predictions according to these information criteria?

*# Both PSIS and WAIC values suggest that model m6.9 (that considers the effect of age and
 # marriage status on happiness) is expected to make better predictions than model m6.10 (that
 # omits marriage status)*

```
# Smaller values of PSIS and WAIC are better.
```

b) On the basis of the causal model, how should you interpret the parameter estimates from the model preferred by PSIS and WAIC?

```
# The pWAIC and pPSIS are the penalty terms.
```

```
# These values are close to the number of dimensions (3 for m6.9 and 2 for m6.10)
```

```
# in the posterior of each model
```

```
# The columns dWAIC and dPSIS reflect the difference between each model's WAIC/PSIS and
```

```
# the best WAIC/PSIS in the set. So it's zero for the best model and then the
```

```
# differences with the other models tell you how far apart each is from the top model.
```

```
# Model m6.9 is about 330 units of deviance smaller than models m6.10.
```

```
# SE is the approximate standard error of each WAIC/PSIS. In a very approximate sense,
```

```
# we expect out-of-sample accuracy to be normally distributed with mean equal to the
```


```
# reported WAIC/PSIS value and a standard deviation equal to the standard error.
```

```
# To judge whether two models are easy to distinguish, we don't use their standard
```

```
# errors but rather the standard error of their difference.
```


```
# visually understanding PSIS based difference in two models
```

```
plot( compare( m6.9 , m6.10, func=PSIS ) )
```



```
# visually understanding WAIC based difference in two models
```

```
plot( compare( m6.9 , m6.10, func=WAIC ) )
```



*# The filled points are the in-sample deviance values. The open points are the WAIC values.
 # Each model does better in-sample than it is expected to do out-ofsample.
 # The line segments show the standard error of each WAIC. These are the values
 # of SE in the table above. So we can see how much better m6.9 is than m6.10.
 # The standard error of the difference in WAIC between the two models is shown by the lighter
 # line segment with the triangle on it, between m6.9 and m6.10.*

2. Laffer Curve

In 2007 The Wall Street Journal published an editorial arguing that raising corporate tax rates increases government revenues only to a point, after which higher tax rates produce less revenue for governments. The editorial included the following graph of corporate tax rates in 29 countries plotted against tax revenue, over which a Laffer curve was drawn. The data used in this plot are available in the rethinking package.

```
library(rethinking)
data(Laffer)
d3 <- Laffer
precis( d3 )[,1:4]
```

```
##           mean      sd   5.5%   94.5%
## tax_rate    26.382069 8.753579 10.9862 35.3568
## tax_revenue  3.306207 1.816491  1.4336  5.7522
```


- a) Using this data, fit a basic regression that uses tax rate to predict tax revenue. Simulate and justify your priors.

```
library(dagitty)
dag_q2 <- dagitty('dag{ taxRATE -> taxREVENUE }')
drawdag( dag_q2 )
```

taxRATE

taxREVENUE

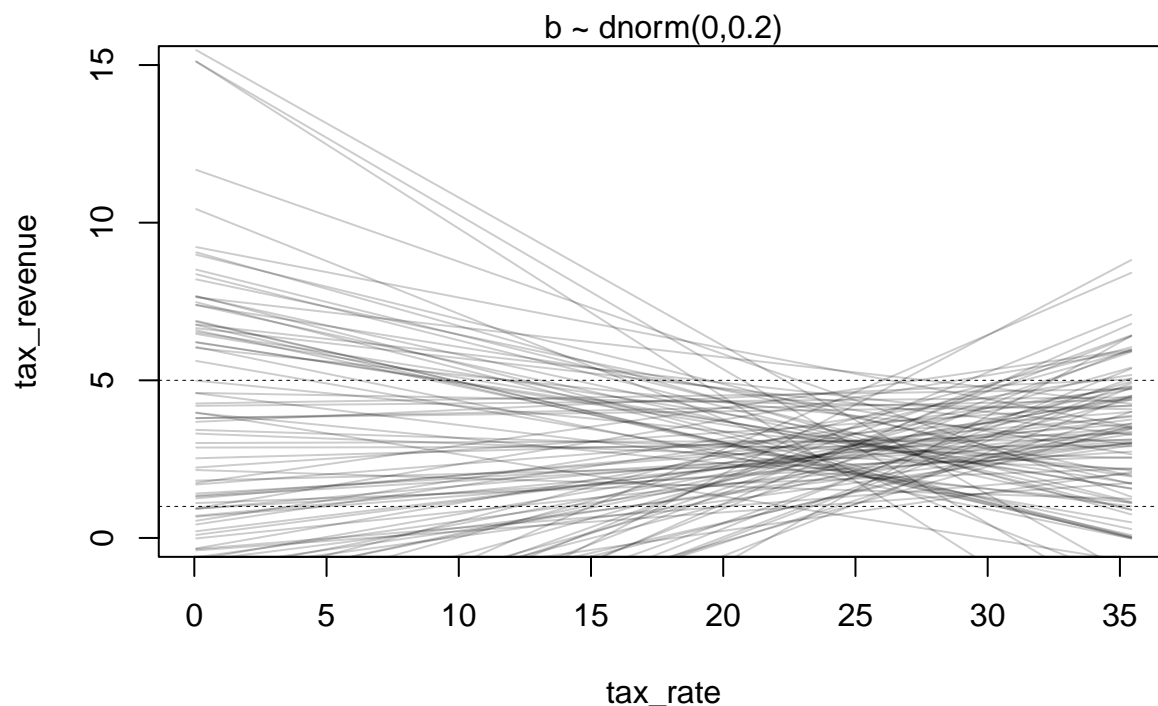
```
# plotting the data  
plot( d3$tax_revenue ~ d3$tax_rate )
```



```
# Presenting prior predictive simulation for tax_revenue & tax_rate model with  
# normal-distribution of 'b'  
# For starters, the intercept 'a' (mean tax_revenue) can assume a normal distribution  
# with a mean of 3  
# While the slope 'b' can be assumed to be normally distributed and centered at mean 0,
```

```
# with SD of 0.2
```

```
N <- 100
a <- rnorm( N , 3 , 1 )
b <- rnorm( N , 0 , 0.2 )
plot( NULL , xlim=range(d3$tax_rate) , ylim=c(0,15) ,
      xlab="tax_rate" , ylab="tax_revenue" )
abline( h=1 , lty=2, lwd=0.5 )
abline( h=5 , lty=2, lwd=0.5 )
mtext( "b ~ dnorm(0,0.2)" )
xbar <- mean(d3$tax_rate)
for ( i in 1:N ) curve( a[i] + b[i]*(x - xbar) ,
                        from=min(d3$tax_rate) , to=max(d3$tax_rate) , add=TRUE ,
                        col=col.alpha("black",0.2) )
```

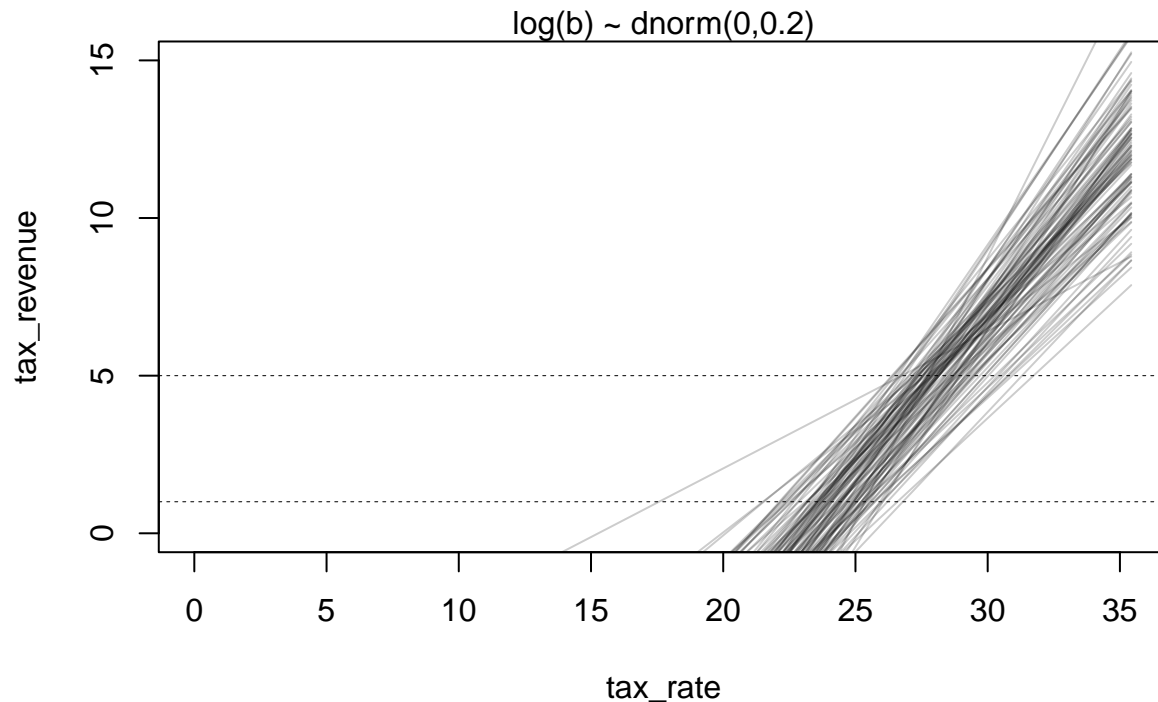


```
# The values of b are completely random in first iteration suggesting tax_revenue could
# be related to tax_rate in any random fashion, which should not be the ideal case.
```

```
# Attempting prior predictive simulation for tax_revenue & tax_rate model with
# log-normal-distribution of 'b'. Let's try restricting it to positive values
# assuming that average tax_revenue increases with average tax_rate, at least up to a point.
```

```
N <- 100
a <- rnorm( N , 3 , 1 )
b <- rlnorm( N , 0 , 0.2 )
# Prior predictive simulation for the tax_revenue and tax_rate model
plot( NULL , xlim=range(d3$tax_rate) , ylim=c(0,15) ,
      xlab="tax_rate" , ylab="tax_revenue" )
abline( h=1 , lty=2, lwd=0.5 )
abline( h=5 , lty=2, lwd=0.5 )
mtext( "log(b) ~ dnorm(0,0.2)" )
xbar <- mean(d3$tax_rate)
for ( i in 1:N ) curve( a[i] + b[i]*(x - xbar) ,
```

```
from=min(d3$tax_rate) , to=max(d3$tax_rate) , add=TRUE ,
col=col.alpha("black",0.2) )
```



```
# Posterior distribution for tax_revenue and tax_rate - basic linear model
# fit model
xbar <- mean(d3$tax_rate)
m2a <- quap(
  alist(
    tax_revenue ~ dnorm( mu , sigma ) ,
    mu <- a + b*( tax_rate - xbar ) ,
    a ~ dnorm( 3 , 1 ) ,
    b ~ dlnorm( 0 , 0.2 ) ,
    sigma ~ dunif( 0 , 4 )
  ) ,
  data=d3 )
# the marginal posterior distributions is as follows
precis( m2a )
```

```
##          mean          sd      5.5%      94.5%
## a      3.2321651 0.49219863 2.4455366 4.0187935
## b      0.3600818 0.05836684 0.2668003 0.4533633
## sigma  3.0411687 0.57780854 2.1177191 3.9646184
```


```
round( vcov( m2a ) , 3 )
```

```
##          a          b      sigma
## a      0.242 -0.001 -0.012
## b     -0.001  0.003  0.024
## sigma -0.012  0.024  0.334
```

```
# the variance-covariance matrix
```



```
plot( tax_revenue ~ tax_rate , data=d3 , col=range(2) )
post <- extract.samples( m2a )
a_map <- mean(post$a)
b_map <- exp(mean(post$b))
curve( a_map + b_map*(x - xbar) , add=TRUE )
```



b) Now construct and fit any curved model you wish to the data. Plot your straight-line model and your new curved model. Each plot should include 89% PI intervals.

```
# Posterior distribution for tax_revenue and tax_rate making use of a
# QUADRATIC_POLYNOMIAL based model
d3$tax_rate_s <- ( d3$tax_rate - mean(d3$tax_rate) )/sd(d3$tax_rate)
d3$tax_rate_s2 <- d3$tax_rate_s^2
m2b_2 <- quap(
  alist(
    tax_revenue ~ dnorm( mu , sigma ) ,
    mu <- a + b1*tax_rate_s + b2*tax_rate_s2 ,
    a ~ dnorm( 3 , 1 ) ,
    b1 ~ dlnorm( 0 , 0.2 ) ,
    b2 ~ dnorm( 0 , 0.2 ) ,
    sigma ~ dunif( 0 , 4 )
  ) ,
  data=d3 )
precis(m2b_2)
```

```
##           mean          sd      5.5%    94.5%
## a      3.32166224 0.3259466  2.8007367  3.8425878
## b1      0.85882695 0.1510481  0.6174230  1.1002309
## b2     -0.04947294 0.1383547 -0.2705904  0.1716445
## sigma   1.70678289 0.2267825  1.3443406  2.0692251
```


```
# Plotting the LINEAR model now
tax_rate.seq <- seq(from=-50, to=50, length.out=50)
```

```

pred_dat <- list( tax_rate=tax_rate.seq )
mu <- link( m2a , data=pred_dat )
mu.mean <- apply( mu , 2 , mean )
mu.PI <- apply( mu , 2 , PI , prob=0.89 )
sim.tax_revenue <- sim( m2a , data=pred_dat )
tax_revenue.PI <- apply( sim.tax_revenue , 2 , PI , prob=0.89 )

plot( d3$tax_revenue ~ d3$tax_rate , d , col=col.alpha(rangi2,0.5) )
lines( tax_rate.seq , mu.mean )
shade( mu.PI , tax_rate.seq )
shade( tax_revenue.PI , tax_rate.seq )

```




```

# Plotting the QUADRATIC POLYNOMIAL model now
tax_rate.seq <- seq(from=-3.5, to=3.5, length.out=50)
pred_dat <- list( tax_rate_s=tax_rate.seq , tax_rate_s2=tax_rate.seq^2 )
mu <- link( m2b_2 , data=pred_dat )
mu.mean <- apply( mu , 2 , mean )
mu.PI <- apply( mu , 2 , PI , prob=0.89 )
sim.tax_revenue <- sim( m2b_2 , data=pred_dat )
tax_revenue.PI <- apply( sim.tax_revenue , 2 , PI , prob=0.89 )

plot( d3$tax_revenue ~ d3$tax_rate_s , d , col=col.alpha(rangi2,0.5) )
lines( tax_rate.seq , mu.mean )
shade( mu.PI , tax_rate.seq )
shade( tax_revenue.PI , tax_rate.seq )

```



c) Using WAIC or PSIS, compare a straight-line model to your curved model. What conclusions would you draw from comparing your two models?

```
# comparing models using WAIC
compare( m2a, m2b_2, func=WAIC)
```

##		WAIC	SE	dWAIC	dSE	pWAIC	weight
##	m2b_2	123.9816	21.41884	0.000	NA	5.819023	9.999995e-01
##	m2a	152.9196	8.24238	28.938	17.61486	3.614630	5.202275e-07

```
# comparing models using PSIS
compare( m2a, m2b_2, func=PSIS)
```

```
## Some Pareto k values are high (>0.5). Set pointwise=TRUE to inspect individual points.
```


```
## Some Pareto k values are very high (>1). Set pointwise=TRUE to inspect individual points.
```

##		PSIS	SE	dPSIS	dSE	pPSIS	weight
##	m2b_2	126.8757	24.333475	0.0000	NA	7.237916	9.999983e-01
##	m2a	153.4591	8.690183	26.5834	19.60764	3.848416	1.688444e-06

```
# Both PSIS and WAIC values suggest that QUADRATIC model m2b_2 is expected to make better
# predictions than LINEAR model m2a
# Smaller values of PSIS and WAIC are better.
# The pWAIC and pPSIS are the penalty terms.
# The columns dWAIC and dPSIS reflect the difference between each model's WAIC/PSIS and
# the best WAIC/PSIS in the set. So it's zero for the best model and then the
# differences with the other models tell you how far apart each is from the top model.
# Model m2b_2 is about 28 units of deviance smaller than models m6.10.
```


```
# SE is the approximate standard error of each WAIC/PSIS. In a very approximate sense,
# we expect out-of-sample accuracy to be normally distributed with mean equal to the
# reported WAIC/PSIS value and a standard deviation equal to the standard error.
# To judge whether two models are easy to distinguish, we don't use their standard
# errors but rather the standard error of their difference.
```

```
plot( compare( m2a, m2b_2, func=WAIC ) )
```



```
plot( compare( m2a, m2b_2, func=PSIS ) )
```

Some Pareto k values are very high (>1). Set pointwise=TRUE to inspect individual points.
 ## Some Pareto k values are very high (>1). Set pointwise=TRUE to inspect individual points.




*# When we look at the plots, we realise that model m2b_2 has a higher SE (~24) as compared to
the SE (~8) of m2a . But m2b_2 is still a favorable choice of model because of its low PSIS and
WAIC values*

- d) There is one country with a high tax revenue which is an outlier. Use PSIS and WAIC to measure the importance of this outlier in the two models you fit.

```
set.seed(24071847)
PSIS_m2a <- PSIS(m2a, pointwise=TRUE)
```

Some Pareto k values are very high (>1). Set pointwise=TRUE to inspect individual points.


```
set.seed(24071847)
WAIC_m2a <- WAIC(m2a, pointwise=TRUE)
plot( PSIS_m2a$k , WAIC_m2a$penalty , xlab="PSIS Pareto k" ,
      ylab="WAIC penalty" , col=range(2) , lwd=2 )
```



```
set.seed(24071847)
PSIS_m2b_2 <- PSIS(m2b_2,pointwise=TRUE)
```

Some Pareto k values are very high (>1). Set pointwise=TRUE to inspect individual points.

```
set.seed(24071847)
WAIC_m2b_2 <- WAIC(m2b_2,pointwise=TRUE)
plot( PSIS_m2b_2$k , WAIC_m2b_2$penalty , xlab="PSIS Pareto k" ,
      ylab="WAIC penalty" , col="rangi2" , lwd=2 )
```




```
# In both of these cases, there is an identifiable outlier with high Pareto k value as well as
# high WAIC penalty.
# We need to account for this outlier because Points like these are highly influential and
# potentially hurt prediction.
```

```
# Let's re-estimate the rate-revenue model model using a Student-t distribution with  $\nu = 4$ 
# we call the model m2b_2t
m2b_2t <- quap(
  alist(
    tax_revenue ~ dstudent( 4 , mu , sigma ) ,
    mu <- a + b1*tax_rate_s + b2*tax_rate_s2 ,
    a ~ dnorm( 3 , 1 ) ,
    b1 ~ dlnorm( 0 , 0.2 ) ,
    b2 ~ dnorm( 0 , 0.2 ) ,
    sigma ~ dunif( 0 , 4 )
  ) , data = d3 )
precis(m2b_2t)
```

```
##           mean      sd      5.5%      94.5%
## a      3.00860735 0.2522217  2.6055083 3.4117064
## b1      0.78652992 0.1326318  0.5745588 0.9985011
## b2     -0.04162556 0.1059588 -0.2109682 0.1277170
## sigma   1.07598878 0.1995687  0.7570394 1.3949382
```

```
# When we compute PSIS now, PSIS(m2b_2t), we don't get any warnings about Pareto k values.
set.seed(24071847)
PSIS_m2b_2t <- PSIS(m2b_2t,pointwise=TRUE)
set.seed(24071847)
WAIC_m2b_2t <- WAIC(m2b_2t,pointwise=TRUE)
plot( PSIS_m2b_2t$k , WAIC_m2b_2t$penalty , xlab="PSIS Pareto k" ,
      ylab="WAIC penalty" , col=rangi2 , lwd=2 )
```




```

# NOW, Plotting the QUADRATIC model m2b_2t (with Student's T distribution)
tax_rate.seq <- seq(from=-3.5, to=3.5, length.out=50)
pred_dat <- list( tax_rate_s=tax_rate.seq , tax_rate_s2=tax_rate.seq^2 )
mu <- link( m2b_2t , data=pred_dat )
mu.mean <- apply( mu , 2 , mean )
mu.PI <- apply( mu , 2 , PI , prob=0.89 )
sim.tax_revenue <- sim( m2b_2t , data=pred_dat )
tax_revenue.PI <- apply( sim.tax_revenue , 2 , PI , prob=0.89 )

plot( d3$tax_revenue ~ d3$tax_rate_s , d , col=col.alpha(rangi2,0.5) )
lines( tax_rate.seq , mu.mean )
shade( mu.PI , tax_rate.seq )
shade( tax_revenue.PI , tax_rate.seq )

```



From the above graph as well, we can see that now the 89% PI intervals for m2b_2t (with Student-t distribution) is much more narrow (confident) than that for the model b2b_2

e) Given your analysis, what conclusions do you draw about the relationship between tax rate and tax revenue? Do your conclusions support the original Laffer curve plot used in the editorial?

*# From the analysis, it is easy to conclude that a polynomial regressor (degree 2) is a better predictor of tax_revenue than a linear regressor (degree 1).
 # This is much in lines with the Laffer curve plot which also somewhat resembles a quadratic polynomial (degree 2)
 # Even though with significantly high PSIS/WAIC values, the relationship between tax_revenue and tax_rate is certainly better explained by a degree two polynomial than a degree one linear curve.*