

# Computational Statistics & Probability

## Problem Set 4 - HMC and Generalized Linear Models

Due: 23:59:59 7.dec.2022

Fall 2022

### Instructions

Assignments must be submitted through Canvas. See the course Canvas page for policies covering collaboration, acceptable file formats (.Rmd & .pdf), and late submissions. Completed assignments must include executable code (.Rmd) and a corresponding knitted markdown file (pdf). An R Markdown [cheat sheet](#) is available.

**Note:** You must have `rstan` installed to complete this assignment.

### 1. Log-odds

a) If an event has probability 0.3, what are the log-odds of this event?

```
p <- 0.3
log( p / (1- p) )
```

```
## [1] -0.8472979
```

b) If an event has log-odds of 1, what is the probability of that event?

```
# use `logistic()`
logistic( 1 )
```

```
## [1] 0.7310586
```

c) If a logistic regression coefficient has value -0.70, what does this imply about the proportional change in odds of the outcome? Briefly explain your answer.

```
# use `exp()`
exp( -0.70 )
```

```
## [1] 0.4965853
```

```
# This means that each unit change in the predictor variable multiplies the odds
# of the event by roughly 1/2, or exp(-0.70) = 0.4965853.
```

### 2. HMC, Interactions and Robust Priors

Recall the interaction model `m8.3`, which is a varying-slope regression model assessing the effect of a country being inside or outside Africa on relationship between the ruggedness of its terrain and its GDP.

```
m8.3 <- quap(
  alist(
    log_gdp_std ~ dnorm(mu, sigma),
    mu <- a[cid] + b[cid]*(rugged_std - 0.215) ,
    a[cid] ~ dnorm(1, 0.1),
    b[cid] ~ dnorm(0, 0.3),
```

```

    sigma ~ dexp(1)
  ), data = dd) # See R code 9.11 to prepare dd

library(rethinking)
data(rugged)
d <- rugged
d$log_gdp <- log(d$rgdppc_2000)
dd <- d[ complete.cases(d$rgdppc_2000) , ]
dd$log_gdp_std <- dd$log_gdp / mean(dd$log_gdp)
dd$rugged_std <- dd$rugged / max(dd$rugged)
dd$cid <- ifelse( dd$cont_africa==1, 1, 2 )
m8.3 <- quap(
  alist(
    log_gdp_std ~ dnorm(mu, sigma),
    mu <- a[cid] + b[cid]*(rugged_std - 0.215) ,
    a[cid] ~ dnorm(1, 0.1),
    b[cid] ~ dnorm(0, 0.3),
    sigma ~ dexp(1)
  ), data = dd)

```

a) Now fit this same model using Hamiltonian Monte Carlo (HMC). The code to do this is in the book, beginning with R code 9.13. You should use the `ulam` convenience function provided by the `rethinking` package.

```

dat_slim <- list(
  log_gdp_std = dd$log_gdp_std,
  rugged_std = dd$rugged_std,
  cid = as.integer( dd$cid )
)

m9.1 <- ulam(
  alist(
    log_gdp_std ~ dnorm(mu, sigma),
    mu <- a[cid] + b[cid]*(rugged_std - 0.215) ,
    a[cid] ~ dnorm(1, 0.1),
    b[cid] ~ dnorm(0, 0.3),
    sigma ~ dexp(1)
  ), data=dat_slim, chains=4)

```

```
show( m9.1 )
```

```

## Hamiltonian Monte Carlo approximation
## 2000 samples from 4 chains
##
## Sampling durations (seconds):
##           warmup sample total
## chain:1   0.06   0.04  0.10
## chain:2   0.06   0.04  0.10
## chain:3   0.07   0.04  0.11
## chain:4   0.06   0.04  0.10
##
## Formula:
## log_gdp_std ~ dnorm(mu, sigma)
## mu <- a[cid] + b[cid] * (rugged_std - 0.215)
## a[cid] ~ dnorm(1, 0.1)

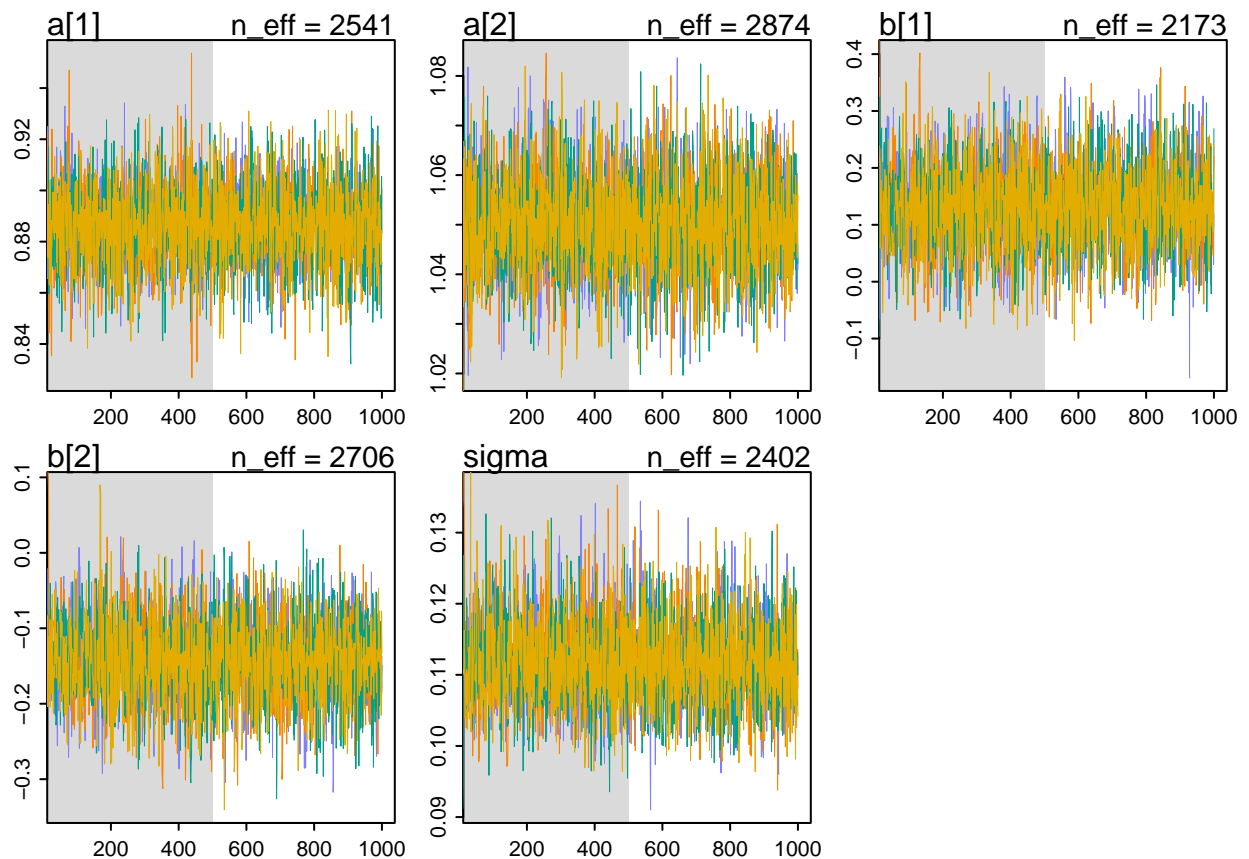
```

```
## b[cid] ~ dnorm(0, 0.3)
## sigma ~ dexp(1)
precis( m9.1, depth=2 )
```

	mean	sd	5.5%	94.5%	n_eff	Rhat4
a[1]	0.8865946	0.015897560	0.86026301	0.91269809	2541.480	0.9985800
a[2]	1.0502974	0.010107657	1.03431945	1.06675110	2874.485	0.9993111
b[1]	0.1328096	0.075558031	0.01112909	0.25882495	2172.861	0.9992068
b[2]	-0.1446571	0.055629771	-0.23601964	-0.05674836	2706.148	0.9995358
sigma	0.1115347	0.006263894	0.10191389	0.12196116	2402.418	0.9992819

b) Check your chains with traceplots and tankplots. Interpret these graphs to explain why, or why not, your HMC model is suitable for inference.

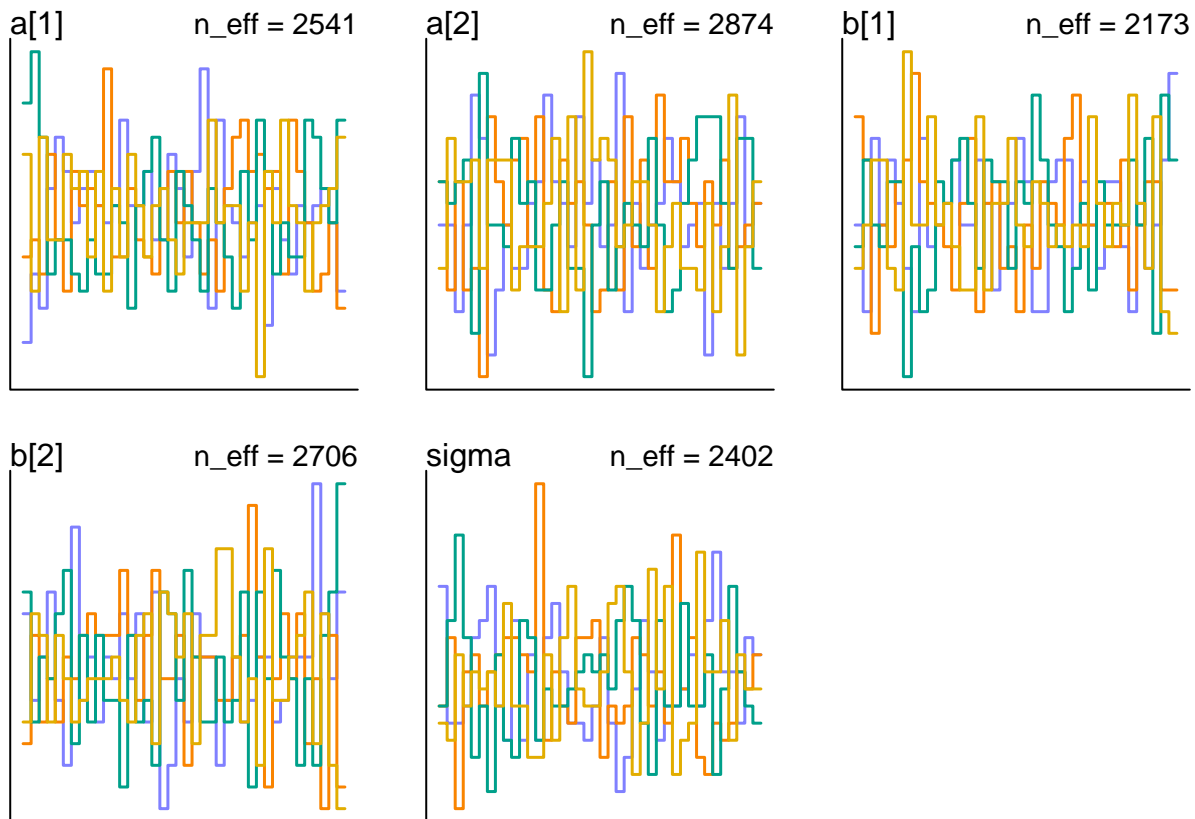
```
traceplot( m9.1, window = c(50,1000))
trankplot( m9.1 )
```



*# DISCUSSION: We include here the trace plot, which a window that omits the first 49 samples with `window = c(50,100)`. We want to check these plots for stationarity, (are the paths of the chains at the same height; that is, is there a stable central tendency that each chain obeys), mixture (do the paths quickly move to explore the full region), and convergence (do multiple, independent chains stick around the same region with high probability).*

*# Next, we include a trace rank (trank) plot, which takes all samples for each parameter and ranks them. When the chains explore the same space efficiently, the histograms are similar to one another and largely overlap.*

```
# Which is what we see. You should always check the health of your chains
# when running HMC.
```



c) Now fit your HMC model with a flat prior for sigma,  $\text{sigma} \sim \text{dunif}(0,1)$ . What effect does this prior have on your posterior distribution? Explain your answer.

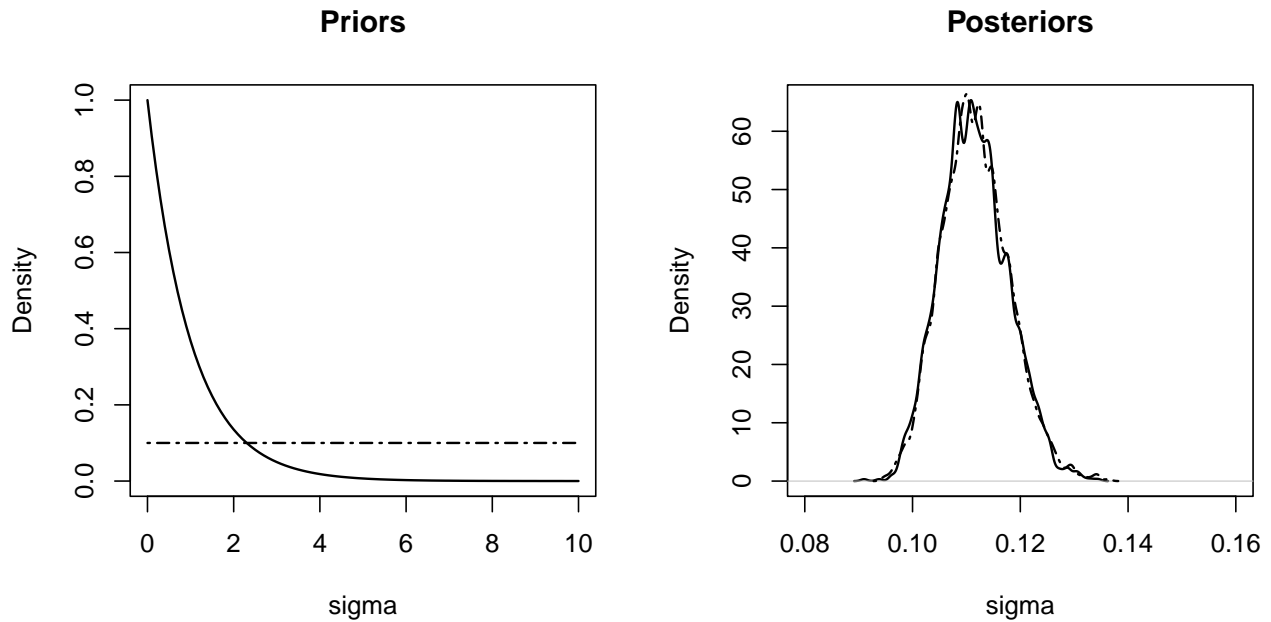
```
m9.1_unif <- ulam(
  alist(
    log_gdp_std ~ dnorm(mu, sigma),
    mu <- a[cid] + b[cid]*(rugged_std - 0.215) ,
    a[cid] ~ dnorm(1, 0.1),
    b[cid] ~ dnorm(0, 0.3),
    sigma ~ dunif(0,1)
  ), data=dat_slim, chains=4)
precis( m9.1_unif )

sigma <- extract.samples(m9.1, pars="sigma")
sigma_unif <- extract.samples(m9.1_unif, pars="sigma")
```

```
# Below is a comparison of prior distributions to posterior distributions for
# sigma. The left plot displays the original model m9.1 with an informed prior
# (solid) and m9.1_unif with an uninformed uniform prior (dashed) on sigma.
```

```
par(mfrow=c(1,2))
# priors
curve( dexp(x,1) , from=0 , to=10 , main="Priors",
      xlab="sigma" , ylab="Density" , ylim=c(0,1), lwd=1.5 )
curve( dunif(x, 0, 10) , add=TRUE, lwd=1.5, lty=6)
```

```
# posterior
dens( sigma[[1]] , xlab="sigma", xlim=c(0.08, 0.16), main="Posteriors", lwd=1.5)
dens( sigma_unif[[1]] , add=TRUE, lwd=1.5, lty=6)
```







*# DISCUSSION: Despite the difference between the informed, regularizing exponential prior (solid) and the very weak uniform prior (dashed), the posterior distributions for sigma are practically identical.*

*# Indeed, we see no effective difference between the posterior distributions:*

*# Plot guide:*

*# TOP Row: Baseline model m9.1, African (blue) and Non-African (red)*

*# Bottom Row: Alternative m9.1\_unif, African (blue) and Non-African (red)*



d) Now fit your model with the log normal prior  $b[cid] \sim \text{dlnorm}(0,1)$  for  $b$ . What effect does this prior have on your posterior distribution? Explain your answer.

```
m9.1_lnorm <- ulam(
  alist(
    log_gdp_std ~ dnorm(mu, sigma),
    mu <- a[cid] + b[cid]*(rugged_std - 0.215) ,
    a[cid] ~ dnorm(1, 0.1),
    b[cid] ~ dlnorm(0, 1),
    sigma ~ dexp(1)
  )
)
```

```

), data=dat_slim, chains=4)

b <- extract.samples(m9.1, pars="b")
b <- b[[1]]
b_lnorm <- extract.samples(m9.1_lnorm, pars="b")
b_lnorm <- b_lnorm[[1]]

# DISCUSSION: The change of the variable slope parameter prior `b` from a normal
# distribution, which allows positive and negative slopes, to a log-normal, which
# rules out negative slopes, has the effect of misfitting a positive slope to
# the non-African countries when it should be negative
precis( m9.1, depth= 2)

##              mean          sd          5.5%          94.5%        n_eff        Rhat4
## a[1]    0.8865946 0.015897560 0.86026301 0.91269809 2541.480 0.9985800
## a[2]    1.0502974 0.010107657 1.03431945 1.06675110 2874.485 0.9993111
## b[1]    0.1328096 0.075558031 0.01112909 0.25882495 2172.861 0.9992068
## b[2]   -0.1446571 0.055629771 -0.23601964 -0.05674836 2706.148 0.9995358
## sigma  0.1115347 0.006263894 0.10191389 0.12196116 2402.418 0.9992819

precis( m9.1_unif, depth=2)

##              mean          sd          5.5%          94.5%        n_eff        Rhat4
## a[1]    0.8862651 0.015795512 0.861259430 0.9115735 2995.977 0.9989617
## a[2]    1.0505513 0.010072209 1.034530000 1.0665966 3299.377 0.9993064
## b[1]    0.1306254 0.077108492 0.009004215 0.2507081 2510.780 0.9994029
## b[2]   -0.1416425 0.057174813 -0.228306445 -0.0526917 2577.316 1.0001769
## sigma  0.1117303 0.006347419 0.102080560 0.1222537 2391.314 0.9993651


# This difference is visualized in terms of the posterior distributions for `b`,
# followed by a comparison of the baseline regression models of m9.1 to those
# of model m9.1_lnorm.

par(mfrow=c(1,2))
dens( b[,2] , xlab="b", xlim=c(-1, 1), main="Posteriors m9.1", col="tomato", lwd=2)
dens( b[,1] , add=TRUE, col="steelblue", lwd=2)
abline(v=0, lty=2, lwd=0.5)


dens( b_lnorm[,2] , xlab="b", xlim=c(-1, 1), main="Posteriors m9.1_lnorm",
      col="tomato", lwd=2, lty=6)
dens( b_lnorm[,1] , add=TRUE, col="steelblue", lwd=2, lty=6)
abline(v=0, lty=2, lwd=0.5)

```





Posteriors m9.1



Posteriors m9.1\_Inorm







### 3. Binomial Regression

We started the course sampling marbles from a bucket to estimate its contents and tossing a globe to estimate the proportion of its surface covered in water. Each made use of the binomial distribution and was ideal to introduce the fundamentals of Bayesian inference. Nevertheless, *Binomial regression* – which is *any* type of GLM using a binomial mean-variance relationship – introduces complications that we needed to postpone until now.

Return to the prosocial chimpanzee experiment in section §11.1 of the textbook, and the HMC model that features individual chimpanzee (**actor**) parameters **actor** and individual **treatment** parameters:

a) Compare m11.4 to a Laplacian quadratic approximate posterior distribution, constructed using `quap()`, that also includes individual parameters for `actor` and `treatment`. What are the differences and similarities between the two approximate posteriors? Explain your answer.

```
data(chimpanzees)
dat <- chimpanzees
dat$treatment <- 1 + dat$prosoc_left + 2*dat$condition
dat_list <- list(
  pulled_left = dat$pulled_left,
  actor = dat$actor,
  treatment = as.integer(dat$treatment) )

# Original HMC Model
m11.4 <- ulam(
  alist(
    pulled_left ~ dbinom( 1, p ),
    logit(p) <- a[actor] + b[treatment] ,
    a[actor] ~ dnorm( 0 , 1.5 ),
    b[treatment] ~ dnorm( 0, 0.5 )
  ), data = dat_list, chains=4, log_lik =TRUE)

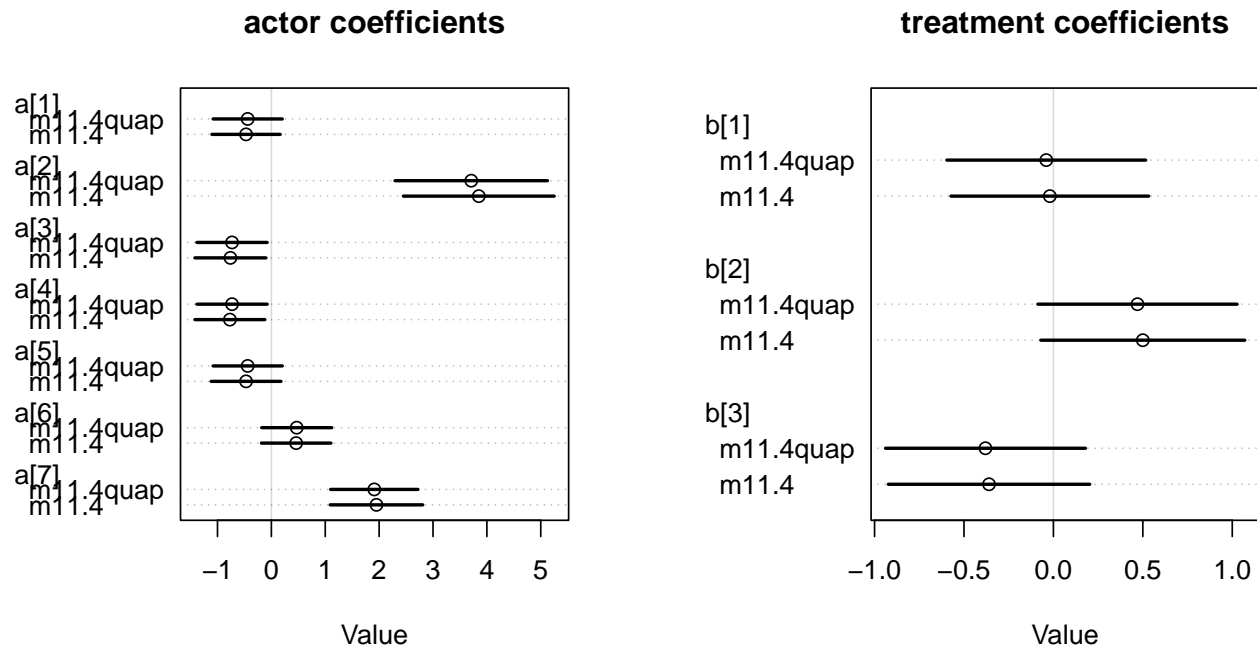
# Quadratic Approximation Model
m11.4quap <- quap(
  alist(
    pulled_left ~ dbinom( 1, p ),
    logit(p) <- a[actor] + b[treatment] ,
    a[actor] ~ dnorm( 0 , 1.5 ),
    b[treatment] ~ dnorm( 0, 0.5 )
  ), data = dat)

#plot( precis( m11.4, depth=2, pars="a"))

# A quick inspection of the parameters of each reveals that the estimates are
# identical, except for actor 2 `a[2]`, where quadratic approximation gives a
# slightly lower estimate of uncertainty on handedness than the HMC model. On
# closer inspection, we see a tendency of all quap estimates of the actors'
# handedness to revert very slightly to the mean. To be clear, these differences
# are dominated by SE for making any strong claim about the parameter values. But,
# this overall pattern does suggest a systematic difference in the numerical
# approximation methods (Laplace vs HMC), a difference we can investigate.

# That difference is most pronounced for actor 2. So let's focus there.

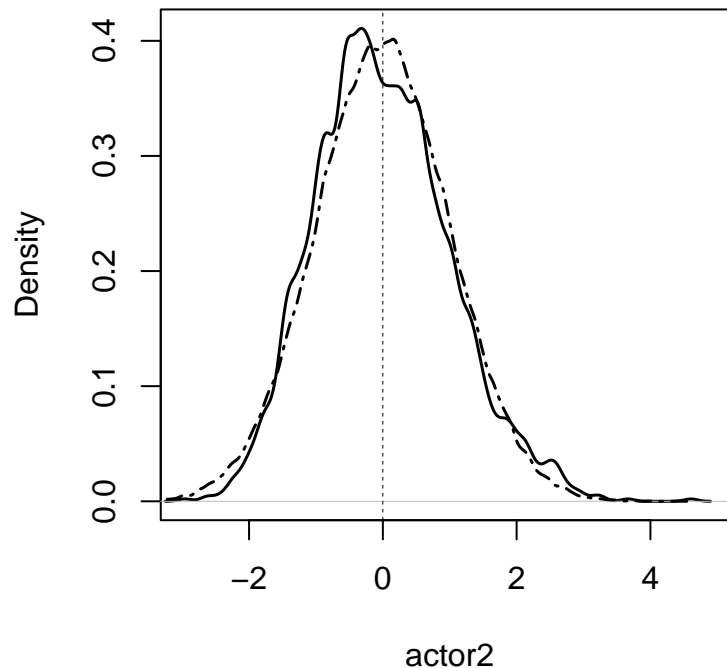
par(mfrow=c(1,2))
plot( coeftab( m11.4, m11.4quap), pars=c("a[1]", "a[2]", "a[3]", "a[4]", "a[5]", "a[6]", "a[7]" ), main
plot( coeftab( m11.4, m11.4quap), pars=c("b[1]", "b[2]", "b[3]" ), main="treatment coefficients")
```



```
# We extract samples from each posterior for the second actor, `a[2]`, and plot
# samples from each posterior distribution, m11.4 (solid) and m11.4quap (dashed)
actor2 <- extract.samples(m11.4, pars="a[2]")
actor2 <- standardize(actor2[[1]])
actor2quap <- extract.samples(m11.4quap, pars="a[2]")
actor2quap <- standardize(actor2quap[[1]])

dens( actor2 , xlab="actor2", xlim=c(-3, 5), main="Posteriors", lwd=1.5)
dens( actor2quap , add=TRUE, lwd=1.5, lty=6)
abline(v=0, lty=2, lwd=0.5)
# We can see the slight difference in each models MAP estimates.
```

## Posteriors



b) Change the prior on the variable intercept to `dnorm( 0 , 10)` and estimate the posterior distribution with both `ulam()` and `quap()`. Do the differences between the two estimations increase, decrease, or stay the same? Explain your answer.

```
# Modified HMC Model
m11.4_b <- ulam(
  alist(
    pulled_left ~ dbinom( 1, p ),
    logit(p) <- a[actor] + b[treatment] ,
    a[actor] ~ dnorm( 0 , 10 ),
    b[treatment] ~ dnorm( 0, 0.5 )
  ), data = dat_list, chains=4, log_lik =TRUE)

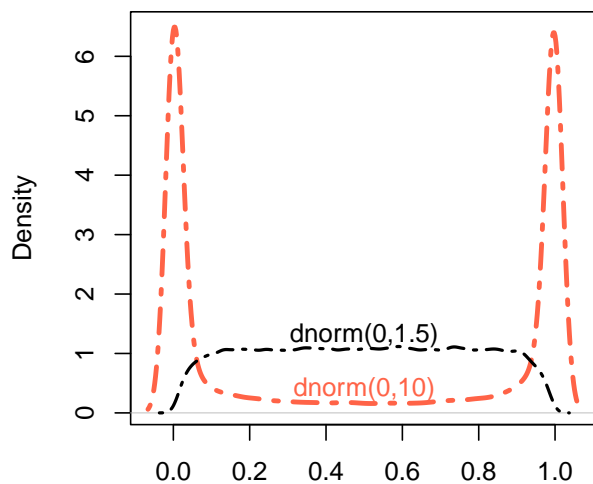
# Modified Quadratic Approximation Model
m11.4quap_b <- quap(
  alist(
    pulled_left ~ dbinom( 1, p ),
    logit(p) <- a[actor] + b[treatment] ,
    a[actor] ~ dnorm( 0 , 10 ),
    b[treatment] ~ dnorm( 0, 0.5 )
  ), data = dat)

# extract priors Laplacian models
prior_crazy <- extract.prior(m11.4quap_b, n=10000)
prior_sane <- extract.prior(m11.4quap, n=10000)
# extract priors HMC models
prior_crazy2 <- extract.prior(m11.4_b, n=10000)
prior_sane2 <- extract.prior(m11.4, n=10000)
```

*# The first thing to recognize is that `dnorm(0, 10)` is a *stronger prior* than `dnorm(0, 1.5)`. The reason is that this "flat prior" is on the parameter scale, # but once you convert it to the outcome scale you see that the prior expresses # that every chimpanzee is expected to always pull the left lever or never pull # the left lever:*


```
par(mfrow=c(1,2))
# Laplacian models
dens(inv_logit(prior_crazy$a), lty=6, lwd=3, col="tomato", main="Laplacian model priors")
dens(inv_logit(prior_sane$a), add=TRUE, lty=6, lwd=2)
text(0.5, 0.4, "dnorm(0,10)", col="tomato")
text(0.5, 1.3, "dnorm(0,1.5)")
# HMC models
dens(inv_logit(prior_crazy2$a), lwd=3, col="tomato", main="HMC priors")
dens(inv_logit(prior_sane2$a), add=TRUE, lwd=2)
text(0.5, 0.4, "dnorm(0,10)", col="tomato")
text(0.5, 1.3, "dnorm(0,1.5)")
```

**Laplacian model priors**



N = 70000 Bandwidth = 0.02219

**HMC priors**




N = 70000 Bandwidth = 0.02219

*# We plot the pair of priors for the Laplacian (quap) model and HMC (ulam) model # simply to show, or remind you, that the priors are the same: neither are mediated # by the model. The crazy prior `dnorm(0,10)` is plotted in red, the sensible # prior `dnorm(0,1.5)` in black.*


*# What effects if any does this strong prior have? A quick inspection of the # parameters of each confirms that the estimates remain identical, except for # actor 2 `a[2]`, where quadratic approximation gives an even lower estimate of # uncertainty on handedness:*

```
par(mfrow=c(1,2))
plot( coeftab( m11.4_b, m11.4quap_b), pars=c("a[1]", "a[2]", "a[3]", "a[4]", "a[5]", "a[6]", "a[7]" ), )
plot( coeftab( m11.4_b, m11.4quap_b), pars=c("b[1]", "b[2]", "b[3]" ), main="treatment coefficients")
```

**actor coefficients**



**treatment coefficients**

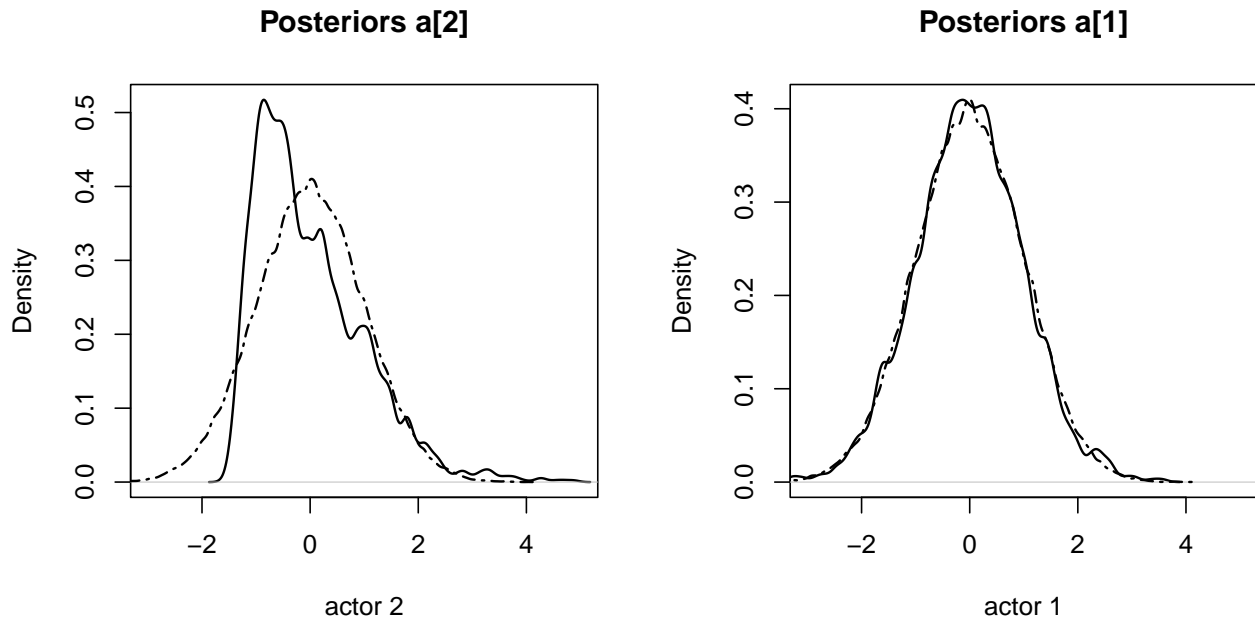


```
# Let us extract samples from each posterior for the second actor, `a[2]`, and plot
# samples from each posterior distribution, m11.4 (solid) and m11.4quap (dashed)
actor2_b <- extract.samples(m11.4_b, pars="a[2]")
actor2_b <- standardize(actor2_b[[1]])
actor2quap_b <- extract.samples(m11.4quap_b, pars="a[2]")
actor2quap_b <- standardize(actor2quap_b[[1]])

# For comparison, consider actor 1
actor1_b <- extract.samples(m11.4_b, pars="a[1]")
actor1_b <- standardize(actor1_b[[1]])
actor1quap_b <- extract.samples(m11.4quap_b, pars="a[1]")
actor1quap_b <- standardize(actor1quap_b[[1]])

# PLOT actor 2
dens( actor2_b , xlab="actor 2", xlim=c(-3, 5), main="Posteriors a[2]", lwd=1.5)
dens( actor2quap_b , add=TRUE, lwd=1.5, lty=6)
#abline(v=0, lty=2, lwd=0.5)

# PLOT actor 1
dens( actor1_b , xlab="actor 1", xlim=c(-3, 5), main="Posteriors a[1]", lwd=1.5)
dens( actor1quap_b , add=TRUE, lwd=1.5, lty=6)
```



*# DISCUSSION: There are a couple things going on here. First, the data dominates the strong prior for all actors except the one,  $a[2]$ , which only pulls the left lever. As there is a lot more mass in the prior  $\text{dnorm}(0,10)$  on this possibility, than the relatively flat (in outcome space) prior  $\text{dnorm}(0,1.5)$ , the ulam variable-effects model responds to this in the posterior.*

*# The strong prior does not have the same effect on the Laplacian varying effects model, however. Why? Quadratic approximation presumes that the posterior distribution is (multivariate) Gaussian, which enforces symmetric uncertainty around the posterior mode on each dimension. But, uncertainty within logistic regression is sometimes asymmetric, which is what we see in the case of handedness of actor 2. The HMC model can respond to the strong prior,  $\text{dnorm}(0,10)$ , pushing the consequences of that strong assumption -- namely that chimps always pull the right or always pull the left, and your uncertainty is simply which hand it will be -- through to your posterior distribution. Actor 2 accords with your strong prior: the evidence suggests he's a Lefty.*

*# What is the effect on the posterior? As values of  $x$  approach infinity, the logistic function returns an estimation of the probability that actor 2 will pull the left lever that approaches 1. The long tail of the m11.4 posterior for  $a[2]$  represents the one-sided infinity of large, positive values that are consistent with the data (i.e., of observing only left-handed pulls from  $N$  pulls). In the small world of m11.4, the strong  $\text{dnorm}(0,10)$  prior simply inflates this tale to emphasize your uncertainty of how close to 1 your estimate should be given  $N$  observations of actor 2. In the large world, the assumption that  $\text{dnorm}(0,10)$  encodes is unreasonable if not crazy.*

*# Lastly, notice that this crazy prior only impacts estimates concerning actor 2's handedness. To see this, compare the asymmetric uncertainty regarding the prospect of actor 2 ever pulling the right handle to the posterior for actor 1, where the uncertainty in the posterior distribution of is approximately Gaussian. Hence, the quadratic approximate posterior of m11.4quap accords with the full Bayesian distribution of m11.4 everywhere else, despite of the crazy  $\text{dnorm}(0,10)$  prior.*